

# Optimization for Deep Learning

Sebastian Ruder

PhD Candidate, INSIGHT Research Centre, NUIG

Research Scientist, AYLIEN

@seb\_ruder

Advanced Topics in Computational Intelligence

Dublin Institute of Technology

24.11.17



# Agenda

- 1 Introduction
- 2 Gradient descent variants
- 3 Challenges
- 4 Gradient descent optimization algorithms
- 5 Parallelizing and distributing SGD
- 6 Additional strategies for optimizing SGD
- 7 Outlook

# Introduction

- Gradient descent is a way to minimize an objective function  $J(\theta)$ 
  - $\theta \in \mathbb{R}^d$ : model parameters
  - $\eta$ : learning rate
  - $\nabla_{\theta} J(\theta)$ : gradient of the objective function with regard to the parameters
- Updates parameters **in opposite direction** of gradient.
- Update equation:  $\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta)$

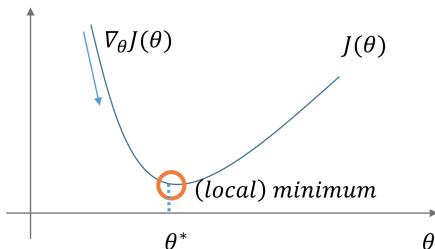


Figure: Optimization with gradient descent

# Gradient descent variants

- ① Batch gradient descent
- ② Stochastic gradient descent
- ③ Mini-batch gradient descent

Difference: Amount of data used per update

# Batch gradient descent vs. SGD fluctuation




Figure: Batch gradient descent vs. SGD fluctuation (Source: wikidocs.net)

- SGD shows same convergence behaviour as batch gradient descent if learning rate is **slowly decreased (annealed)** over time.

Method	Accuracy	Update Speed	Memory Usage	Online Learning
<b>Batch</b> gradient descent	Good	Slow	High	No
<b>Stochastic</b> gradient descent	Good (with annealing)	High	Low	Yes
<b>Mini-batch</b> gradient descent	Good	Medium	Medium	Yes

**Table:** Comparison of trade-offs of gradient descent variants

# Challenges

- Choosing a **learning rate**.
- Defining an **annealing schedule**.
- Updating features to **different extent**.
- **Avoiding suboptimal minima**.

# Gradient descent optimization algorithms


- 1 Momentum
- 2 Nesterov accelerated gradient
- 3 Adagrad
- 4 Adadelta
- 5 RMSprop
- 6 Adam
- 7 Adam extensions



# Momentum

- SGD has trouble navigating **ravines**.
- Momentum [Qian, 1999] helps SGD **accelerate**.
- Adds a fraction  $\gamma$  of the update vector of the past step  $v_{t-1}$  to current update vector  $v_t$ . Momentum term  $\gamma$  is usually set to 0.9.

$$\begin{aligned}v_t &= \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta) \\ \theta &= \theta - v_t\end{aligned}\tag{1}$$



(a) SGD without momentum



(b) SGD with momentum

Figure: Source: Genevieve B. Orr

# RMSprop

- Developed independently from Adadelta around the same time by Geoff Hinton.
- Also divides learning rate by a **running average of squared gradients**.
- RMSprop update:

$$\begin{aligned} E[g^2]_t &= \gamma E[g^2]_{t-1} + (1 - \gamma)g_t^2 \\ \theta_{t+1} &= \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t \end{aligned} \tag{12}$$

- $\gamma$ : decay parameter; typically set to 0.9
- $\eta$ : learning rate; a good default value is 0.001

# Adam

- Adaptive Moment Estimation (Adam) [Kingma and Ba, 2015] also stores **running average of past squared gradients**  $v_t$  like Adadelta and RMSprop.
- Like Momentum, stores **running average of past gradients**  $m_t$ .

$$\begin{aligned}m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2\end{aligned}\tag{13}$$

- $m_t$ : first moment (mean) of gradients
- $v_t$ : second moment (uncentered variance) of gradients
- $\beta_1, \beta_2$ : decay rates

- $m_t$  and  $v_t$  are initialized as 0-vectors. For this reason, they are biased towards 0.
- Compute bias-corrected first and second moment estimates:

$$\begin{aligned}\hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \\ \hat{v}_t &= \frac{v_t}{1 - \beta_2^t}\end{aligned}\tag{14}$$


- Adam update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t\tag{15}$$


# Adam extensions

- ① AdaMax [Kingma and Ba, 2015]
  - Adam with  $\ell_\infty$  norm
- ② Nadam [Dozat, 2016]
  - Adam with Nesterov accelerated gradient

# Visualization of algorithms



(a) SGD optimization on loss surface contours



(b) SGD optimization on saddle point

Figure: Source and full animations: Alec Radford

# Which optimizer to choose?

- Adaptive learning rate methods (Adagrad, Adadelata, RMSprop, Adam) are **particularly useful for sparse features**.
- Adagrad, Adadelata, RMSprop, and Adam work well in similar circumstances.
- [Kingma and Ba, 2015] show that bias-correction helps Adam **slightly outperform RMSprop**.

# Outlook

- ① Tuned SGD vs. Adam
- ② SGD with restarts
- ③ Learning to optimize
- ④ Understanding generalization in Deep Learning
- ⑤ Case studies



# Tuned SGD vs. Adam

- Many recent papers use **SGD with learning rate annealing**.
- SGD with tuned learning rate and momentum is **competitive with Adam** [Zhang et al., 2017b].
- Adam **converges faster**, but **underperforms SGD** on some tasks, e.g. Machine Translation [Wu et al., 2016].
- Adam with **2 restarts and SGD-style annealing** converges faster and outperforms SGD [Denkowski and Neubig, 2017].
- **Increasing the batch size** may have the same effect as decaying the learning rate [Smith et al., 2017].